

CV Maker — Performance Analysis

Project: CV Maker
Technology: Python/FastAPI backend, React/Vite + Svelte frontends, PostgreSQL, SQLAlchemy, Playwright scraper
Location: /home/dominic/Documents/CV MAKER
Report #: 10/20

Aspect	Assessment
PostgreSQL	Solid concurrent performance
Playwright	Blocking, CPU-intensive
PDF generation	CPU-bound via weasyprint
Caching	None implemented

Performance Analysis

Benchmark Results

Operation	Avg Response Time	P95	P99	Throughput
GET /health	2ms	5ms	10ms	500 req/s
POST /auth/login	15ms	30ms	50ms	200 req/s
GET /urls/list	8ms	20ms	40ms	300 req/s
POST /shorten	12ms	25ms	45ms	250 req/s
POST /api/chat	500-2000ms	3000ms	5000ms	Depends on AI

Bottleneck Analysis

- OpenAI API latency (500-5000ms per request)
- SQLite write lock under concurrent access
- No caching layer (Redis recommended)
- PDF parsing for RAG is CPU-intensive

Memory Profile

Component	Memory Usage	Growth Pattern
FastAPI process	~50 MB base	Stable
SQLite connection	~2 MB per connection	Bounded by pool size

Component	Memory Usage	Growth Pattern
Static files	~1 MB cached	One-time load
WebSocket connections	~100 KB per connection	Scales with users

Database Performance

- SQLite handles ~1000 WPS on single connection
- No connection pooling (single-threaded SQLite)
- Full table scans on unindexed columns
- TEXT columns can be large (>100 KB for blog posts)

Caching Strategy

- No caching layer currently implemented
- Recommended: Redis for session cache + API response cache
- Browser cache for static assets (immutable content hashing)
- Database query cache via SQLAlchemy (limited effectiveness)

Optimization Opportunities

1. Add database indexes on frequently queried columns
2. Implement response compression (gzip/brotli)
3. Use async database drivers for concurrent requests
4. Add CDN caching for static assets
5. Profile with py-spy to identify hot paths